

Test of problems from SuiteSparse Matrix Collection

This package provides implementations of comparing BICGSTAB, GMRES, DQGMRES, FOM, DIOM, SCG, SWI, RGMRES, RFOM, RSCG, and DSWI for solving the linear system

$$Ax = b,$$

where the matrix A from the SuiteSparse Matrix Collection and set b so that the solution is $x_\star = (1, 1, \dots, 1)$.

Contents

- BICGSTAB
- REGMRES
- DQGMRES
- REFOM
- DIOM(m)
- RESCG
- SWI(m)
- DYNOSWIWP
- Test

BICGSTAB

Using Matlabs function “bicgstab” directly.

REGMRES

REGMRES (Restarted GMRES): Algorithm 6.11 in Yousef Saad’s “Iterative Methods for Sparse Linear System (2nd Edition)”

```
function [x, k, res, resvec] = regmres(A, b, restart, tol, x0, maxit)
```

- $[x] = \text{regmres}(A, b)$ attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .
- $[x] = \text{regmres}(A, b, \text{restart})$ specifies the restarted number. If restart is $[]$ then regmres uses the default, n
- $[x] = \text{regmres}(A, b, \text{restart}, \text{tol})$ specifies the tolerance of the method. If tol is $[]$ then regmres uses the default, $1e-6$.
- $[x] = \text{regmres}(A, b, \text{restart}, \text{tol}, x0)$ specifies the initial guess. If $x0$ is $[]$ then regmres uses the default, an all zero vector.
- $[x] = \text{regmres}(A, b, \text{restart}, \text{tol}, x0, \text{maxit})$ specifies the maximum number of iterations. If maxit is $[]$ then regmres uses the default, 10000.

- `[x, k] = regmres(A, b, ...)` returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$.
- `[x, k, res] = regmres(A, b, ...)` also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- `[x, k, res, resvec] = regmres(A, b, ...)` also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

GMRES: set `restart = n` in the function “`regmres`”.

DQGMRES

DQGMRES (Direct Quasi-GMRES): Algorithms 6.6 and 6.13 in Yousef Saad’s “Iterative Methods for Sparse Linear System (2nd Edition)”

```
function [x, k, res, resvec] = dqgmres(A, b, m, tol, x0, maxit)
```

- `[x] = dqgmres(A, b)` attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .
- `[x] = dqgmres(A, b, m)` specifies the number of the sliding window. If `m` is `[]` then `dqgmres` uses the default, `n`.
- `[x] = dqgmres(A, b, m, tol)` specifies the tolerance of the method. If `tol` is `[]` then `dqgmres` uses the default, `1e-6`.
- `[x] = dqgmres(A, b, m, tol, x0)` specifies the initial guess. If `x0` is `[]` then `dqgmres` uses the default, an all zero vector.
- `[x] = dqgmres(A, b, m, tol, x0, maxit)` specifies the maximum number of iterations. If `maxit` is `[]` then `dqgmres` uses the default, `10000`.
- `[x, k] = dqgmres(A, b, ...)` returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$.
- `[x, k, res] = dqgmres(A, b, ...)` also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- `[x, k, res, resvec] = dqgmres(A, b, ...)` also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

REFOM

REFOM (Restarted FOM): Algorithm 6.5 in Yousef Saad’s “Iterative Methods for Sparse Linear System (2nd Edition)”

```
function [x, k, res, resvec] = refom(A, b, restart, tol, x0, maxit)
```

- `[x] = refom(A, b)` attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .

- $[x] = \text{refom}(A, b, \text{restart})$ specifies the restarted number. If restart is $[]$ then refom uses the default, n
- $[x] = \text{refom}(A, b, \text{restart}, \text{tol})$ specifies the tolerance of the method. If tol is $[]$ then refom uses the default, 1e-6.
- $[x] = \text{refom}(A, b, \text{restart}, \text{tol}, x0)$ specifies the initial guess. If x0 is $[]$ then refom uses the default, an all zero vector.
- $[x] = \text{refom}(A, b, \text{restart}, \text{tol}, x0, \text{maxit})$ specifies the maximum number of iterations. If maxit is $[]$ then refom uses the default, 10000.
- $[x, k] = \text{refom}(A, b, \dots)$ returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$.
- $[x, k, \text{res}] = \text{refom}(A, b, \dots)$ also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- $[x, k, \text{res}, \text{resvec}] = \text{refom}(A, b, \dots)$ also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

FOM: set restart= n in the function “refom”.

DIOM(m)

DIOM (Direct Incomplete Orthogonalization Method): Algorithms 6.6 and 6.8 in Yousef Saad’s “Iterative Methods for Sparse Linear System (2nd Edition)”

```
function [x, k, res, resvec] = diom(A, b, mk, tol, x0, maxit)
```

- $[x] = \text{diom}(A, b)$ attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .
- $[x] = \text{diom}(A, b, m)$ specifies the number of the sliding window. If m is $[]$ then diom uses the default, n.
- $[x] = \text{diom}(A, b, m, \text{tol})$ specifies the tolerance of the method. If tol is $[]$ then diom uses the default, 1e-6.
- $[x] = \text{diom}(A, b, m, \text{tol}, x0)$ specifies the initial guess. If x0 is $[]$ then diom uses the default, an all zero vector.
- $[x] = \text{diom}(A, b, m, \text{tol}, x0, \text{maxit})$ specifies the maximum number of iterations. If maxit is $[]$ then diom uses the default, 10000.
- $[x, k] = \text{diom}(A, b, \dots)$ returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$.
- $[x, k, \text{res}] = \text{diom}(A, b, \dots)$ also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- $[x, k, \text{res}, \text{resvec}] = \text{diom}(A, b, \dots)$ also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

RESCG

RESCG (Restarted SCG): Restarted semi-conjugate gradient method

```
function [x, k, res, resvec] = rescg(A, b, restart, tol, x0, maxit)
```

- $[x] = \text{rescg}(A, b)$ attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .
- $[x] = \text{rescg}(A, b, \text{restart})$ specifies the restarted number. If restart is $[]$ then rescg uses the default, n .
- $[x] = \text{rescg}(A, b, \text{restart}, \text{tol})$ specifies the tolerance of the method. If tol is $[]$ then rescg uses the default, $1e-6$.
- $[x] = \text{rescg}(A, b, \text{restart}, \text{tol}, x0)$ specifies the initial guess. If $x0$ is $[]$ then rescg uses the default, an all zero vector.
- $[x] = \text{rescg}(A, b, \text{restart}, \text{tol}, x0, \text{maxit})$ specifies the maximum number of iterations. If maxit is $[]$ then rescg uses the default, 10000.
- $[x, k] = \text{rescg}(A, b, \dots)$ returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$.
- $[x, k, \text{res}] = \text{rescg}(A, b, \dots)$ also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- $[x, k, \text{res}, \text{resvec}] = \text{rescg}(A, b, \dots)$ also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

SCG: set $\text{restart}=n$ in the function “ rescg ”.

SWI(m)

SWI: Sliding window implementation with pre-allocated memory

```
[x, k, res, resvec] = swi(A, b, m, tol, x0, maxit)
```

- $[x] = \text{swi}(A, b)$ attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .
- $[x] = \text{swi}(A, b, m)$ specifies the number of the sliding window. If m is $[]$ then swi uses the default, n .
- $[x] = \text{swi}(A, b, m, \text{tol})$ specifies the tolerance of the method. If tol is $[]$ then swi uses the default, $1e-6$.
- $[x] = \text{swi}(A, b, m, \text{tol}, x0)$ specifies the initial guess. If $x0$ is $[]$ then swi uses the default, an all zero vector.
- $[x] = \text{swi}(A, b, m, \text{tol}, x0, \text{maxit})$ specifies the maximum number of iterations. If maxit is $[]$ then swi uses the default, 10000.
- $[x, k] = \text{swi}(A, b, \dots)$ returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$.
- $[x, k, \text{res}] = \text{swi}(A, b, \dots)$ also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- $[x, k, \text{res}, \text{resvec}] = \text{swi}(A, b, \dots)$ also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

DYNSWIWP

DYNSWIWP: Sliding window implementation with choosing m dynamically

The approach to choose m dynamically is as follows:

$$m = \begin{cases} \max\{[0.8m], 2\}, & \text{if } \|r_k\| < 0.9\|r_{k-1}\|; \\ \min\{2m, n\}, & \text{if } \|r_k\| > 2\|r_{k-1}\|; \\ m, & \text{otherwise.} \end{cases}$$

```
[x, k, res, resvec] = dynswiwp(A, b, m, tol, x0, maxit)
```

- $[x] = \text{dynswiwp}(A, b)$ attempts to find a solution x to the system of linear equations $Ax = b$. The n -by- n coefficient matrix A must be positive definite but need not be symmetric. The right hand side column vector b must have length n .
- $[x] = \text{dynswiwp}(A, b, m)$ specifies the number of the sliding window. If m is $[]$ then `dynswiwp` uses the default, n .
- $[x] = \text{dynswiwp}(A, b, m, \text{tol})$ specifies the tolerance of the method. If `tol` is $[]$ then `dynswiwp` uses the default, $1e-6$.
- $[x] = \text{dynswiwp}(A, b, m, \text{tol}, x0)$ specifies the initial guess. If $x0$ is $[]$ then `dynswiwp` uses the default, an all zero vector.
- $[x] = \text{dynswiwp}(A, b, m, \text{tol}, x0, \text{maxit})$ specifies the maximum number of iterations. If `maxit` is $[]$ then `dynswiwp` uses the default, 10000.
- $[x, k] = \text{dynswiwp}(A, b, \dots)$ returns the iteration number at which x was computed: $1 \leq k \leq \text{maxit}$
- $[x, k, \text{res}] = \text{dynswiwp}(A, b, \dots)$ also returns the last relative residual norm $\|b - Ax\|/\|b\|$.
- $[x, k, \text{res}, \text{resvec}] = \text{dynswiwp}(A, b, \dots)$ also returns a vector of estimates of the residual norms at each iteration, including $\|b - Ax\|$.

Test

1. Download the files: `diom.m`, `dqgmres.m`, `swi.m`, `swiwp.m`, `cage12.mat`, `test_suitesparse.m`, and `testproblem.txt`;
2. Run the function “`test_suitesparse`” directly, i.e.,

```
>> test_suitesparse
```
3. The test results are saved in “`TestResult`” file.